

Chapter IV - Adversarial Search

4.1 Games

Games are a common and useful domain for testing artificial intelligence techniques. In AI, a game is often modeled as a search problem where agents (players) take turns making decisions to maximize their chances of winning.

Types of Games:

- **Deterministic Games:** Games with no element of chance (e.g., Chess, Tic-Tac-Toe).
- **Stochastic Games:** Games that involve randomness (e.g., Poker, Backgammon).
- **Perfect Information Games:** Games where all players have full knowledge of the game state (e.g., Chess, Checkers).
- **Imperfect Information Games:** Games where players have incomplete knowledge (e.g., Poker).

Game Representation:

- **Initial State:** Describes the starting configuration.
 - **Players:** Defines who plays the game (e.g., MAX and MIN in two-player games).
 - **Actions:** The possible moves each player can make.
 - **Transition Model:** Defines how moves change the game state.
 - **Terminal State:** Defines when the game ends and assigns a value (win, lose, draw).
 - **Utility Function:** A numerical representation of game outcomes (e.g., +1 for win, 0 for draw, -1 for loss).
-

4.2 Optimal Decisions in Games

To make optimal decisions in adversarial environments, AI agents must evaluate possible future moves and choose the best one. The **Minimax Algorithm** is a fundamental method used in perfect information games.

Minimax Algorithm:

- **MAX Player:** Tries to maximize the score.
- **MIN Player:** Tries to minimize the score (acting as an opponent).

- **Tree Representation:** The game is represented as a tree where nodes represent possible game states, and edges represent moves.
- **Backtracking:** The algorithm explores all possible outcomes and propagates values back to the root to determine the best move.

Example:

- In a Tic-Tac-Toe game, the minimax algorithm simulates every possible move and selects the one leading to the best outcome.

Limitations:

- **Computationally Expensive:** In games with large branching factors (like Chess or Go), minimax can become impractical.
 - **Cannot Handle Uncertainty:** It assumes a fully deterministic environment.
-

4.3 Alpha-Beta Pruning

To improve minimax efficiency, **Alpha-Beta Pruning** eliminates branches in the game tree that cannot possibly influence the final decision.

Key Concepts:

- **Alpha (α):** The best value MAX can achieve so far.
- **Beta (β):** The best value MIN can achieve so far.
- **Pruning:** If a node's evaluation proves that it won't be selected, further exploration is unnecessary, reducing computations.

Effectiveness:

- Reduces the number of nodes evaluated, improving efficiency.
- In an ideal case, it reduces the complexity from $O(b^d)$ to $O(b^{d/2})$, where b is the branching factor and d is the depth.

Example:

- In Chess, if one move is found to be significantly worse than another, we stop evaluating that branch early.
-

4.4 Imperfect Decisions in Real-Time

In real-world scenarios, AI does not have unlimited time to search for the perfect move. Instead, it must make **imperfect but reasonable decisions** under time constraints.

Techniques for Real-Time Decision Making:

1. **Depth-Limited Search:** Search is stopped at a specific depth rather than reaching terminal states.
2. **Heuristic Evaluation Functions:** Instead of computing exact outcomes, heuristics estimate the value of a game state.
3. **Iterative Deepening:** A combination of depth-first and breadth-first search that allows deeper analysis as time permits.
4. **Monte Carlo Tree Search (MCTS):** Simulates multiple possible game playouts to guide decision-making (widely used in Go and modern game AI).

Example:

- In real-time strategy games like StarCraft, AI must make quick decisions based on incomplete information.
-

4.5 Stochastic Games

Stochastic (or probabilistic) games introduce elements of chance, requiring AI to handle randomness in decision-making.

Examples:

- **Dice-based Games:** Backgammon, Monopoly.
- **Card Games:** Poker, Blackjack.

Solution Approaches:

1. **Expectimax Algorithm:** A variation of Minimax that considers probabilistic outcomes. Instead of choosing the maximum or minimum value, it computes the expected value.
 2. **Monte Carlo Methods:** AI simulates thousands of random game scenarios to estimate the best move.
 3. **Reinforcement Learning:** AI learns optimal strategies by playing many games and adjusting based on outcomes.
-

4.6 Partially Observable Games

In **partially observable games**, players do not have full knowledge of the game state. They must infer missing information based on observations and probabilities.

Examples:

- **Poker**: Players do not know opponents' hands.
- **Battleship**: Players only receive limited feedback about opponent actions.

Approaches to Handling Partial Observability:

1. **Belief State Representation**: Instead of a single state, AI maintains a probability distribution over multiple possible states.
 2. **Hidden Markov Models (HMMs)**: Used for tracking hidden information in dynamic systems.
 3. **Bayesian Networks**: Helps AI make decisions under uncertainty by incorporating probabilities.
-

4.7 State-of-the-Art Game Programs

AI has achieved remarkable success in various competitive games:

1. Chess (Deep Blue, 1997)

- IBM's **Deep Blue** defeated world champion Garry Kasparov using brute-force search and evaluation functions.

2. Go (AlphaGo, 2016)

- **AlphaGo** defeated human grandmasters using deep neural networks and reinforcement learning.

3. Poker (Libratus, 2017)

- **Libratus** outperformed professional poker players using game-theoretic reasoning and self-play.

4. StarCraft II (AlphaStar, 2019)

- **AlphaStar** reached **Grandmaster level** using deep reinforcement learning.

5. Dota 2 (OpenAI Five, 2018)

- **OpenAI Five** competed against professional players using reinforcement learning and self-play.
-

4.8 Alternative Approaches

New AI techniques continue to evolve beyond classical search and reinforcement learning.

1. Deep Learning-Based Agents

- Use **convolutional neural networks (CNNs)** and **transformers** to process game data and predict actions.
- Example: **AlphaZero** trained entirely through self-play without human input.

2. Neuroevolution

- Uses **genetic algorithms** to evolve AI agents.
- Example: **NEAT (NeuroEvolution of Augmenting Topologies)** optimizes neural network structures for better decision-making.

3. Hybrid Approaches

- Combines **rule-based systems** with **machine learning** for flexible strategies.
 - Example: AI for **real-time strategy games** often blends symbolic reasoning with deep learning.
-

Exercises

1. **Minimax Algorithm:** Implement a minimax-based Tic-Tac-Toe AI. How does its performance change with increasing depth?
2. **Alpha-Beta Pruning:** Compare the execution time of a minimax algorithm with and without alpha-beta pruning.
3. **Expectimax:** Implement an Expectimax AI for a simple dice-based game (e.g., Pig). How does it handle randomness?
4. **Monte Carlo Tree Search (MCTS):** Implement a basic MCTS algorithm for Connect Four. How does the number of simulations affect performance?
5. **Partially Observable Games:** Design an AI for a simple card game where it must infer hidden cards from observed moves.